Project 3: s2fs (Super Simple File System)

- Handed out: Wednesday, October 16, 2019
- Due dates
 - Parts 1 and 2: Monday, November 25, 2019
 - Parts 3 and 4: Monday, December 9, 2019

Introduction

The goal of this project is to implement a simple virtual file system that can be used to access the information stored in task_struct of each process.

Recommended Background Reading

- Process Management Chapter3 Linux Kernel Development, Robert Love
- File System Implementation
- Creating Linux virtual filesystems
- lwnfs code

Part 1. Print the Process Tree of Linux

Part 1.1. Implement a kernel module

[5 points] Linux kernel stores the information of a process in task_struct structure. A process can create zero or more processes called child process. Children of same process are called sibilings. task_struct stores this information using children and sibling pointer in task struct. The children and sibling pointer together form a tree call the process tree of the kernel. In part 1, you have to iterate recursively from the root of the process tree to print the process tree in a hierarchical fashion (children of same parent should be at same indent level). Please submit you code as a tarball named s2fs_part_1.tar.gz and it should include a makefile to compile and load the kernel module.

Part 1.2. Test your code

[5 points] To test you code, load the kernel module and check the dmesg output. Compare your dmesg output with the output of pstree command. Upload a screen shot of the dmesg output.

```
1081.201547] Loading proctree Module...
1081.201548] swapper/0 [0]
1081.2015491
                systemd [1]
1081.201549]
                  systemd-journal [421]
1081.201550]
                  systemd-udevd [449]
1081.201550]
                  lvmetad [453]
1081.201551]
                  systemd-timesyn [507]
1081.201551
                  systemd-network [649]
1081.201552]
                  systemd-resolve [650]
1081.201552
                  accounts-daemon [774]
                  cron [802]
1081.201553]
1081.201553]
                  irqbalance [803]
1081.201554]
                  systemd-logind [809]
1081.201554]
                  lxcfs [813]
                  atd [818]
1081.201555]
                  snapd [824]
1081.201556]
1081.201556]
                  dbus-daemon [825]
1081.201557]
                  networkd-dispat [934]
1081.201557]
                  rsyslogd [937]
1081.201558]
                  sshd [953]
1081.201558]
                    sshd [1295]
                      sshd [1436]
1081.201559]
1081.201559]
                        bash [1437]
                           sudo [5747]
1081.201560]
                             insmod [5748]
1081.201560]
1081.201561
                  iscsid [988]
1081.201561
                  iscsid [992]
1081.201562
                  polkitd [1065]
                  agetty [1120]
1081.201563]
1081.201563
                  systemd [1303]
1081.201564
                    (sd-pam) [1316]
1081.201564]
                kthreadd [2]
1081.201565]
                  rcu_gp[3]
                  rcu par gp [4]
1081.201566]
                  kworker/0:0H [6]
1081.201566]
1081.201567]
                  mm percpu wq [8]
```

Part 2: Mount and Unmount a Pseudo File system

A pseudo filesystem is not backed by disk but resides in memory and is usually used to provide information regarding the kernel to the user (e.g., proc file system). In Part 2, you will have to create a minimal mountable prototype of a pseduo filesystem which we will name s2fs. For this part, s2fs should mount and unmount on a mount point safely. Implementation of file operation is not required. Please read the following article to understand how to write a simple VFS using libfs.

• Creating Linux virtual filesystems

The article is for Linux 2.6, so some of the API might have been renamed/removed but it is a good starting point. Code that compiles for linux 4.18. Understanding the full source code is highly recommended.

• lwnfs code

Part 2.1: Minimal Mountable Filesystem Prototype

[10 points] Following is a guide to make a minimal muntable filesystem prototype - Writing a File System in Linux Kernel

Please note the above tutorial is for disk backed filesystem. Use the code in Part 2.1 and the above guide to write a filesystem prototype. You do not need to implement file operations. The file system should mount and unmount safely. Submit you code as a tarball named <code>s2fs_part2.tar.gz</code>. Include a makefile to compile, load and mount the kernel module. The filesystem will be mounted to directory mnt in the current folder.

Part 2.2: Test Your Code

[10 points] To check whether you filesystem mounts successfully use the following commands.

```
$ mkdir mnt
$ sudo insmod s2fs.ko # Inserting the filesystem kernel module
$ sudo mount -t s2fs nodev mnt # Mount at dir mnt with option nodev
$ mount # Print out all mounted file systems
$ sudo umount ./mnt # Unmount the file system
$ sudo rmmod s2fs.ko # Remove the kernel module
```

```
tija@chotu:-/project/solution/part25 mkdir mnt
tija@chotu:-/project/solution/part25 ls
Makefile mnt modules.order Module.symvers vtfs.c vtfs.ko vtfs.mod.c vtfs.mod.o vtfs.o
tija@chotu:-/project/solution/part25 sudo insmod vtfs.ko
tija@chotu:-/project/solution/part25 sudo numt - t vtfs.modev mnt
tija@chotu:-/project/solution/part25 sudo mount - t vtfs.modev mnt
tija@chotu:-/project/solution/part25 mount | tail -10
mqueue on /dev/mqueue type mqueue (ms.relatine)
debugfs on /sys/kernel/debug type debugfs (rw.relatine)
mqueue on /dev/mqueue type mqueue (ms.relatine)
mqueue on /dev/mqueue type mqueue (ms.relatine)
mqueue on /dev/mqueque type mqueue (ms.relatine)
mqueue on /dev/mqueque type mqueue (ms.relatine)
mqueue on /dev/mqueue on /dev/mqueue type mqueue on /dev/mqueue type mqueue on /dev/mqueue type mqueue on /
```

Add the screenshot of the output of mount command. If mount is successful, the command will show s2fs as a mounted file system.

Part 3: Implement File Operation

Now we will add inode and file operations to s2fs. In this part, on mounting, s2fs will create a directory named foo and inside the directory it will create a file named bar. You can use code from Part 2.1.

Part 3.1: Function to create Inode with given mode

[5 points] Write a function: static struct inode *s2fs_make_inode(struct super_block *sb, int mode);. The function accepts two input, the superblock of the filesystem and the mode which decides if the inode is for a directory or file and the permission bits. Make sure you set the i_ino field of inode to get_next_ino(). Post a file name and line number of your implementation.

Part 3.2: Function to create directory

[5 points] First write a function to create a directory. The function prototype will be static struct dentry *s2fs_create_dir(struct super_block *sb, struct dentry *parent, const char *dir_name); Post a file name and line number of your implementation.

Part 3.3: Functions to perform file operations

[5 points] To handle a file, the filesystem needs to know how to open, read and write a file. Write three functions, s2fs_open, s2fs_read_file, and s2fs_write_file. In this project we will not use the open and the write fuction so it will just return 0. The read function should return "Hello World!" string to the user. Create a s2fs_fops of type file_operations and assign

.read, .write and .open to the functions you have written. Post a file name and line number where you define s2fs_fops.

Part 3.4: Fuction to create file

[5 points] Write a function: static struct dentry *s2fs_create_file(struct super_block *sb, struct dentry *dir, const char *file_name) It should create a file with name stored in file_name inside directory pointed by dir dentry. Post a file name and line number of your implementation.

Part 3.5: Putting it all together

[5 points] After mounting add code which will create a directory named foo in the root folder and inside foo create a file named bar. They can be created using the s2fs_create_dir and s2fs_create_file functions. Submit a tarball with your code and makefile named s2fs_part3.tar.gz.

Part 3.6: Testing the code

[10 points] First mount the filesystem. The change directory to foo. Then use cat bar to check if reading the file outputs Hello World!

```
$ sudo insmod s2fs.ko
$ sudo mount -t s2fs nodev mnt
$ cd mnt/foo
$ cat bar
$ cd ../..
$ sudo umount ./mnt
$ sudo rmmod s2fs.ko
```

Add the screenshot of the above commands to your submission.

```
tija@chotu:~/project/solution/part2$ make mount
sudo insmod vtfs.ko
sudo mount -t vtfs nodev mnt
tija@chotu:~/project/solution/part2$ cd mnt/
tija@chotu:~/project/solution/part2/mnt$ ls
foo
tija@chotu:~/project/solution/part2/mnt$ cd foo/
tija@chotu:~/project/solution/part2/mnt/foo$ ls
bar
tija@chotu:~/project/solution/part2/mnt/foo$ cat bar
Hello World!
tija@chotu:~/project/solution/part2/mnt/foo$
```

Part 4: Use s2fs to get task information

In part 4, we will use part 1 and part 3 to display information related to task.

Part 4.1: Creating files/directory corresponding to each PID

Part 4.1.1: Create files/directory on mounting

[5 points] Modify the part 3 code to create directories and files at mount time in a recursive manner as explained below:

- If a task has children
 - Create a folder with name as the pid of the task
 - Inside the created folder create a file with the pid of the task
 - Recursively follow the rule for all children of the task
- If a task has no children create a file in it's parent task's folder

For example, if 0 has children 1, 2, 3 and 1 has children 4. Then the directory structure will look like as follows.

0 |--0 |--1 | |--1 | |--4 |--2 |--3

Post a file name and line number of your implementation.

Part 4.1.2: Testing Code

[5 points] Mount the file system and run tree command on the mount folder to see the directory tree. Add screen the screenshot of the same in your submission.

Part 4.2: Reading a file with name X should give task info of PID X

For part 4.2 you have to modify the s2fs_read_file and s2fs_open function, so that when you read a file with name X it will output the task info of task with PID X.

Part 4.2.1: Get Task Info

[5 points] Write a function int get_task_info(int pid, char* data) which takes input a pid and a buffer and returns number of bytes written to buffer and

the task information in the data buffer. If the task does not exists, it should return -1. The information to be returned are:

- Task Name
- Task State
- Process Id
- CPU ID
- Thead Group ID (TGID)
- Parent's PID (PPID)
- Start Time
- Dynamic Priority
- Static Priority
- Normal Priority
- Real-time Priority
- Memory Map Base
- Virtual Memory Space
- Virtual Memory Usage
- No. of Virtual Memory Address
- Total Pages Mapped

Post a file name and line number of your implementation.

Part 4.2.2: Modify the inode creation file to store the PID

[5 points] When creating file for a task, store the PID in inode->i_private. Post a file name and line number where you assign inode->i_private.

Part 4.2.3: Modify the file open and read function

[5 points] Modify the s2fs_open function to get the pid from the file inode and store it in filp->private_data. Modify the s2fs_read_file function to call get_task_info and output the task information instead of Hello World. If the task does not exists, output Task no longer exists. Create a tarball named s2fs_part4.tar.gz of the code and makefile.

Part 4.2.4: Test your code

[10 points] Add screenshot of following commands.

```
$ mkdir mnt
$ sudo insmod s2fs.ko # Inserting the filesystem kernel module
$ sudo mount -t s2fs nodev mnt # Mount at dir mnt with option nodev
$ mount # List all mounted file system
$ cat mnt/0/1/1 # Print out task information of the systemd process
$ sudo umount ./mnt # Unmount the file system
$ sudo rmmod s2fs.ko # Remove the kernel module
$ dmesg
```